

# YICHUN HU

2 W Loop Rd, New York, NY 10044  
yh767@cornell.edu || yichunhu.github.io

## EDUCATION

---

### Cornell University

Ph.D. Candidate in Operations Research  
Ph.D. minor in Applied Mathematics

08/2017 - 05/2023 (*expected*)

### Peking University

B.S. in Mathematics and Applied Mathematics  
B.A. in Economics

09/2013 - 06/2017

## RESEARCH INTERESTS

---

Data-driven Decision-Making, Online/Sequential Learning, Contextual Bandit & Reinforcement Learning, Stochastic Optimization.

## RESEARCH PAPERS

---

### Journal Publications

1. **Y. Hu**, N. Kallus, X. Mao (2022) Smooth Contextual Bandits: Bridging the Parametric and Non-differentiable Regret Regimes.
  - *Operations Research (Articles in Advance)*.
  - **Finalist, INFORMS Applied Probability Society 2020 Best Student Paper Competition**.
  - Preliminary version appeared in 33rd Conference on Learning Theory (*COLT 2020*).
2. **Y. Hu**, N. Kallus, X. Mao (2022) Fast Rates for Contextual Linear Optimization.
  - *Management Science* 68(6):4236-4245.
  - Accepted to Fast Track in Management Science.

### Journal Papers Under Review/Revision

3. **Y. Hu**, N. Kallus. DTR Bandit: Learning to Make Response-Adaptive Decisions with Low Regret.
  - *Under 2nd-round review after major revision at Journal of the American Statistical Association*.
4. **Y. Hu**, N. Kallus, M. Uehara. Fast Rates for the Regret of Offline Reinforcement Learning.
  - Preliminary version appeared in 34th Conference on Learning Theory (*COLT 2021*).
  - Journal version under review.
  - Invited for presentation at Online RL Theory Seminar.

### Conference Papers

5. M. Garrard, H. Wang, B. Letham, Z. Wang, Y. Huang, **Y. Hu**, C. Zhou, N. Zhou, E. Bakshy. Practical Policy Optimization with Personalized Experimentation.
  - *NeurIPS Workshop on Causal Inference Challenges in Sequential Decision Making, 2021*.

## SELECTED HONORS

---

Finalist, Applied Probability Society Best Student Paper Competition, INFORMS	2020
Sherri Koenig Stuewer Graduate Fellowship, Cornell University	2018
Excellent Graduate Award, Peking University	2017
Award for Academic Excellents, Peking University	2015, 2016
May Fourth Scholarship, Peking University	2016
Kwung-Hua Scholarship, Peking University	2015

## INDUSTRY EXPERIENCE

---

### Facebook (Meta)

*Menlo Park, CA (Remote)*

Research Engineer Intern, Core Data Science (Adaptive Experimentation)

*05/2021-08/2021*

- Researched on multi-objective contextual bandit learning and value model tuning in personalized experiments. Resulted in a NeurIPS workshop paper.
- Proposed algorithm has been launched in Facebook's personalized experimentation platform.

### Google

*Mountain View, CA (Remote)*

Data Scientist Intern, Google Play

*05/2020-08/2020*

- Researched on causal methods to analyze the impact of app usage on the retention rate of Google Play Pass.
- Wrote a practitioner-friendly tech report for response curve analysis and treatment effect estimation.

## TEACHING EXPERIENCE

---

### Cornell University, Teaching Assistant

CS 5785: Applied Machine Learning

*Fall 2019*

ORIE 4360: A Mathematical Examination of Fair Representation

*Fall 2018*

ORIE 3510: Introduction to Engineering Stochastic Processes I

*Spring 2018*

ORIE 5600: Financial Engineering with Stochastic Calculus I

*Fall 2017*

## PROFESSIONAL SERVICE

---

- Journal Reviewer: Operations Research, Journal of Machine Learning Research, Transactions on Machine Learning Research
- Conference Reviewer & PC Member: ICML 2020-2022, ICLR 2021, AISTATS 2021-2022, NeurIPS 2021-2022, EAAMO 2022, NeurIPS 2021 Workshop on Causal Inference Challenges in Sequential Decision Making
- Session Chair: INFORMS 2020 General Session (Stochastic Bandits)
- Cornell University ORGA (Operations Research Graduate Association) Tech Liaison *2019-2020*

## SELECTED PRESENTATIONS

---

### Fast Rates for Contextual Linear Optimization

INFORMS Annual Meeting, Indianapolis, IN

*10/2022*

Cornell ORIE Young Researchers Workshop, Ithaca, NY

*10/2022*

NYC Operations Day (Poster), New York, NY

*04/2022*

ORIE PhD Colloquium at Cornell Tech, New York, NY

*04/2022*

INFORMS Optimization Society Conference, Greenville, SC

*03/2022*

Cornell ORIE Young Researchers Workshop (Poster), Ithaca, NY

*10/2021*

### Fast Rates for the Regret of Offline Reinforcement Learning

RL Theory Seminar, Virtual

*11/2021*

INFORMS Annual Meeting, Anaheim, CA

*10/2021*

16th INFORMS Workshop on Data Mining and Decision Analytics, Anaheim, CA

*10/2021*

34th Annual Conference on Learning Theory (COLT 2021), Boulder, CO

*08/2021*

### DTR Bandit: Learning to Make Response-Adaptive Decisions with Low Regret

INFORMS Annual Meeting, Virtual

*11/2020*

15th INFORMS Workshop on Data Mining and Decision Analytics, Virtual

*11/2020*

### Smooth Contextual Bandits: Bridging the Parametric and Non-differentiable Regret Regimes

33rd Annual Conference on Learning Theory (COLT 2020), Virtual

*07/2020*

INFORMS Annual Meeting, Seattle, WA

*10/2019*

14th INFORMS Workshop on Data Mining and Decision Analytics, Seattle, WA

*10/2019*

Cornell ORIE Young Researchers Workshop, Ithaca, NY

*10/2019*

## OTHERS

---

- *Programming:* Python (PyTorch), R, C/C++, Julia.
- *Interests:* Books, Music, Nature, Soccer